

# Big Data Lab | Assignment 5 | Madhur Jindal | ME18B059

1. PFA the screenshot to location of the data file and the code for task 3 in the bucket.

The screenshot shows the Google Cloud Storage interface for a bucket named 'big-data-lab-madhurj'. The bucket is located in 'us (multiple regions in United States)', has a 'Standard' storage class, 'Not public' access, and 'None' protection. The 'OBJECTS' tab is selected, showing a list of objects. The breadcrumb path is 'Buckets > big-data-lab-madhurj > Assignment 5'. The objects listed are 'assignment5\_ME18B059.py' (2.5 KB, text/x-python) and 'iris.data' (4.4 KB, application/octet-stream). Both are stored in the 'Standard' class, last modified on 12 Mar 2022, and are not public. Both use 'Google-managed key' encryption and have no retention or hold policies.

<input type="checkbox"/>	Name	Size	Type	Created	Storage class	Last modified	Public access	Version history	Encryption	Retention expiry date	Holds
<input type="checkbox"/>	assignment5_ME18B059.py	2.5 KB	text/x-python	12 Mar 2022, 21:...	Standard	12 Mar 2022, 21:5...	Not public	—	Google-managed key	—	None
<input type="checkbox"/>	iris.data	4.4 KB	application/octet-stream	12 Mar 2022, 18:...	Standard	12 Mar 2022, 18:4...	Not public	—	Google-managed key	—	None

2. PFA the screenshot to the query for task 2 along with the answer as 0. Hence there do not exist any flowers of the given configuration - Iris Virginica flowers which have sepal width greater than 3 cm and petal length smaller than 2 cm.

## Big Data Lab | Assignment 5 | Madhur Jindal | ME18B059

ASSIGN... X

IRIS\_DATA X

RUN

SAVE

SHARE

SCHEDULE

MORE

```
1 SELECT
2   COUNT(class)
3 FROM
4   `big-data-lab-madhur.iris_data.iris_data`
5 WHERE
6   sepal_width > 3
7   AND petal_length < 2
8   AND class='Iris-virginica'
```

Query results

SAVE RESULTS

EXPLORE DATA

Query complete (0.0 sec elapsed, cached)

Job information

Results

JSON

Row	f0_
1	0

3. The code for the Task 3 is given in the zip named as assignment5\_ME18B059.py file. We try to use different classification models along with different pre-processing techniques on the given IRIS dataset and try to predict the class of the iris plant.

- We start with importing the required libraries and setting up the Spark Context for connection with spark and also setting up the Spark Session using sc for setting up the connection with sql engine.

## Big Data Lab | Assignment 5 | Madhur Jindal | ME18B059

- b. Then we read the table using spark sql from our saved tables as an rdd.
- c. Then we transform the data in the class column to encoded format and also use Vector.dense to transform the other columns to a features column. Finally, we create a Spark DataFrame from the rdd after applying the transform.
- d. We use randomSplit method to create a 4:1 train test split and use MultiClassClassificationEvaluator for calculating accuracy.
- e. Following that we try training different models with no pre-processing on the raw data and also using Standard Scaler for transformation.
- f. The different models that were trained are:
  - a. Logistic Regression
  - b. Decision Tree Classifier
  - c. Random Forest Classifier
  - d. Naïve Bayes Classifier
- g. The training and test accuracies for the combinations have been provided.

```
from __future__ import print_function
from pyspark.context import SparkContext      Import "pyspark.context" could not be resolved
from pyspark.sql.session import SparkSession  Import "pyspark.sql.session" could not be resolved
from pyspark.ml.linalg import Vectors        Import "pyspark.ml.linalg" could not be resolved
from pyspark.ml.classification import LogisticRegression, DecisionTreeClassifier, RandomForestClassifier, NaiveBayes
from pyspark.ml.evaluation import MulticlassClassificationEvaluator      Import "pyspark.ml.evaluation" could not be resolved
from pyspark.ml.feature import StandardScaler      Import "pyspark.ml.feature" could not be resolved

sc = SparkContext()
spark = SparkSession(sc)

data = spark.read.format("bigquery").option("table", "iris_data.iris_data").load()
# data.createOrReplaceTempView("iris")

def vector_from_inputs(r):
    label = 0
    if r['class'] == 'Iris-virginica':
        label = 1
    elif r['class'] == 'Iris-versicolor':
        label = 2
    return (label, Vectors.dense(float(r['sepal_length']),
                                float(r['sepal_width']),
                                float(r['petal_length']),
                                float(r['petal_width'])))

df = data.rdd.map(vector_from_inputs).toDF(['label', 'features'])
```

```
# No Pre-processing

print("-----")
print("NO PREPROCESSING - RAW DATA \n")

df_no_pre = df.select('*')
train, test = df_no_pre.randomSplit([0.8, 0.2])
train.cache()

# Logistic Regression
lr = LogisticRegression(fitIntercept=True, maxIter=10, regParam=0.3)
lrModel = lr.fit(train)

lr_preds_tr = lrModel.transform(train)
lr_preds_tst = lrModel.transform(test)
evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
lr_acc_tr = evaluator.evaluate(lr_preds_tr)
lr_acc_tst = evaluator.evaluate(lr_preds_tst)

print("Logistic Regression :")
print("Train Acc: ", lr_acc_tr)
print("Test Acc: ", lr_acc_tst)
print("\n")

# Decision Tree Classifier

dt = DecisionTreeClassifier()
dtModel = dt.fit(train)

dt_preds_tr = dtModel.transform(train)
dt_preds_tst = dtModel.transform(test)
dt_acc_tr = evaluator.evaluate(dt_preds_tr)
dt_acc_tst = evaluator.evaluate(dt_preds_tst)

print("Decision Tree Classifier :")
print("Train Acc: ", dt_acc_tr)
print("Test Acc: ", dt_acc_tst)
print("\n")
```

```

# Random Forest Classifier

rf = RandomForestClassifier(numTrees=5, seed=42)
rfModel = rf.fit(train)

rf_preds_tr = rfModel.transform(train)
rf_preds_tst = rfModel.transform(test)
rf_acc_tr = evaluator.evaluate(rf_preds_tr)
rf_acc_tst = evaluator.evaluate(rf_preds_tst)

print("Random Forest Classifier :")
print("Train Acc: ", rf_acc_tr)
print("Test Acc: ", rf_acc_tst)
print("\n")

# Naive Bayes Classifier

nb = NaiveBayes(modelType='gaussian')
nbModel = nb.fit(train)

nb_preds_tr = nbModel.transform(train)
nb_preds_tst = nbModel.transform(test)
nb_acc_tr = evaluator.evaluate(nb_preds_tr)
nb_acc_tst = evaluator.evaluate(nb_preds_tst)

print("Naive Bayes Classifier :")
print("Train Acc: ", nb_acc_tr)
print("Test Acc: ", nb_acc_tst)
print("\n")

## Using Standard Scaler

print("-----")
print("STANDARD SCALER \n")

df_standard = df.select('*')
scaler = StandardScaler(inputCol="features", outputCol='Scaledfeatures')
scalerModel = scaler.fit(df_standard)
scaledData = scalerModel.transform(df_standard)
df_scld = scaledData.select('label', 'Scaledfeatures')
df_scaled = df_scld.withColumnRenamed("Scaledfeatures", "features")

train, test = df_scaled.randomSplit([0.8, 0.2])
train.cache()

```

## Big Data Lab | Assignment 5 | Madhur Jindal | ME18B059

```
# Logistic Regression
lr = LogisticRegression(fitIntercept=True, maxIter=10, regParam=0.3)
lrModel = lr.fit(train)

lr_preds_tr = lrModel.transform(train)
lr_preds_tst = lrModel.transform(test)
evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
lr_acc_tr = evaluator.evaluate(lr_preds_tr)
lr_acc_tst = evaluator.evaluate(lr_preds_tst)

print("Logistic Regression :")
print("Train Acc: ", lr_acc_tr)
print("Test Acc: ", lr_acc_tst)
print("\n")

# Decision Tree Classifier
dt = DecisionTreeClassifier()
dtModel = dt.fit(train)

dt_preds_tr = dtModel.transform(train)
dt_preds_tst = dtModel.transform(test)
dt_acc_tr = evaluator.evaluate(dt_preds_tr)
dt_acc_tst = evaluator.evaluate(dt_preds_tst)

print("Decision Tree Classifier :")
print("Train Acc: ", dt_acc_tr)
print("Test Acc: ", dt_acc_tst)
print("\n")

# Random Forest Classifier
rf = RandomForestClassifier(numTrees=5, seed=42)
rfModel = rf.fit(train)

rf_preds_tr = rfModel.transform(train)
rf_preds_tst = rfModel.transform(test)
rf_acc_tr = evaluator.evaluate(rf_preds_tr)
rf_acc_tst = evaluator.evaluate(rf_preds_tst)

print("Random Forest Classifier :")
print("Train Acc: ", rf_acc_tr)
print("Test Acc: ", rf_acc_tst)
print("\n")
```

```
# Naive Bayes Classifier

nb = NaiveBayes(modelType='gaussian')
nbModel = nb.fit(train)

nb_preds_tr = nbModel.transform(train)
nb_preds_tst = nbModel.transform(test)
nb_acc_tr = evaluator.evaluate(nb_preds_tr)
nb_acc_tst = evaluator.evaluate(nb_preds_tst)

print("Naive Bayes Classifier :")
print("Train Acc: ", nb_acc_tr)
print("Test Acc: ", nb_acc_tst)
print("\n")
```



# Big Data Lab | Assignment 5 | Madhur Jindal | ME18B059

-----  
NO PREPROCESSING - RAW DATA

```
22/03/12 17:16:32 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeSystemBLAS
22/03/12 17:16:32 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeRefBLAS
22/03/12 17:16:32 INFO breeze.optimize.StrongWolfeLineSearch: Line search t: 0.15536695335880768 fval: 1.0264374469044402 rhs: 1.0969994104655274 cdd: 0.05257130443998145
22/03/12 17:16:32 INFO breeze.optimize.LBFGS: Step Size: 0.1554
22/03/12 17:16:32 INFO breeze.optimize.LBFGS: Val and Grad Norm: 1.02644 (rel: 0.0643) 1.98523
22/03/12 17:16:32 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:32 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.925262 (rel: 0.0986) 1.48689
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.722709 (rel: 0.219) 0.514207
22/03/12 17:16:33 INFO breeze.optimize.StrongWolfeLineSearch: Line search t: 0.25067292345910097 fval: 0.7156975951187798 rhs: 0.7227076684670068 cdd: -4.315707231231342E-4
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Step Size: 0.2507
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.715698 (rel: 0.00970) 0.124271
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.713316 (rel: 0.00333) 0.0710168
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.711410 (rel: 0.00267) 0.0595227
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:33 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.707175 (rel: 0.00595) 0.0410482
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.705315 (rel: 0.00263) 0.0376701
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.703600 (rel: 0.00243) 0.0625497
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.701175 (rel: 0.00345) 0.0576869
22/03/12 17:16:34 INFO breeze.optimize.LBFGS: Converged because max iterations reached
Logistic Regression :
Train Acc: 0.8245614035087719
Test Acc: 0.7222222222222222
```

```
Decision Tree Classifier :
Train Acc: 1.0
Test Acc: 0.9444444444444444
```

```
Random Forest Classifier :
Train Acc: 1.0
Test Acc: 0.9444444444444444
```

```
Naive Bayes Classifier :
Train Acc: 0.956140350877193
```

Output is complete



# Big Data Lab | Assignment 5 | Madhur Jindal | ME18B059

-----  
STANDARD SCALER

```
22/03/12 17:16:51 INFO breeze.optimize.StrongWolfeLineSearch: Line search t: 0.16371710308137855 fval: 1.0244249536104135 rhs: 1.0983780937404728 cdd: 0.02193584543034889
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 0.1637
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 1.02442 (rel: 0.0673) 2.04202
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.916843 (rel: 0.105) 1.52452
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.734336 (rel: 0.199) 0.974979
22/03/12 17:16:51 INFO breeze.optimize.StrongWolfeLineSearch: Line search t: 0.2738319908124558 fval: 0.7161255621534455 rhs: 0.7343328939802141 cdd: -0.002279651491949207
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 0.2738
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.716126 (rel: 0.0246) 0.0935787
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.715654 (rel: 0.000658) 0.0809403
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.713701 (rel: 0.00273) 0.0743330
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:51 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.709865 (rel: 0.00537) 0.109512
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.708598 (rel: 0.00179) 0.0602954
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.707881 (rel: 0.00101) 0.0358136
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Step Size: 1.000
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Val and Grad Norm: 0.707610 (rel: 0.000384) 0.0605916
22/03/12 17:16:52 INFO breeze.optimize.LBFGS: Converged because max iterations reached
Logistic Regression :
Train Acc: 0.9145299145299145
Test Acc: 0.9696969696969697
```

```
Decision Tree Classifier :
Train Acc: 0.9914529914529915
Test Acc: 0.9696969696969697
```

```
Random Forest Classifier :
Train Acc: 0.9829059829059829
Test Acc: 0.9393939393939394
```

```
Naive Bayes Classifier :
Train Acc: 0.9487179487179487
Test Acc: 0.9696969696969697
```