

Lab7-Subscriber

April 4, 2022

```
[1]: import os
os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages org.apache.spark:
    ↳spark-sql-kafka-0-10_2.12:3.1.2 pyspark-shell'

[2]: from pyspark.sql.types import StructType, StringType, FloatType
from pyspark.sql import SparkSession
from pyspark.sql.functions import explode
from pyspark.sql.functions import split
from pyspark.ml.classification import RandomForestClassificationModel
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.linalg import Vectors
from pyspark.sql import Row
import pyspark.sql.functions as f
from pyspark.ml import PipelineModel
from itertools import chain

spark = SparkSession\
    .builder\
    .appName("Iris-Prediction")\
    .config("spark.driver.extraClassPath", "/home/ubuntu/jars/"
    ↳spark-sql-kafka-0-10_2.12-3.1.2.jar,/home/ubuntu/jars/commons-pool2-2.11.0.
    ↳jar")\
    .getOrCreate()

spark.sparkContext.setLogLevel('WARN')
```

```
:: loading settings :: url = jar:file:/usr/lib/spark/jars/ivy-2.4.0.jar!/org/apache/ivy/core/settings/ivysettings.xml
```

```
Ivy Default Cache set to: /root/.ivy2/cache
```

```
The jars for the packages stored in: /root/.ivy2/jars
```

```
org.apache.spark#spark-sql-kafka-0-10_2.12 added as a dependency
```

```
:: resolving dependencies :: org.apache.spark#spark-submit-
```

```
parent-b4e4886c-b385-40e0-8941-cde23bf18d8d;1.0
```

```
    confs: [default]
```

```
    found org.apache.spark#spark-sql-kafka-0-10_2.12;3.1.2 in central
```

```
    found org.apache.spark#spark-token-provider-kafka-0-10_2.12;3.1.2 in
central
```

```

found org.apache.kafka#kafka-clients;2.6.0 in central
found com.github.luben#zstd-jni;1.4.8-1 in central
found org.lz4#lz4-java;1.7.1 in central
found org.xerial.snappy#snappy-java;1.1.8.2 in central
found org.slf4j#slf4j-api;1.7.30 in central
found org.spark-project.spark#unused;1.0.0 in central
found org.apache.commons#commons-pool2;2.6.2 in central
:: resolution report :: resolve 650ms :: artifacts dl 12ms
:: modules in use:
com.github.luben#zstd-jni;1.4.8-1 from central in [default]
org.apache.commons#commons-pool2;2.6.2 from central in [default]
org.apache.kafka#kafka-clients;2.6.0 from central in [default]
org.apache.spark#spark-sql-kafka-0-10_2.12;3.1.2 from central in
[default]
org.apache.spark#spark-token-provider-kafka-0-10_2.12;3.1.2 from central
in [default]
org.lz4#lz4-java;1.7.1 from central in [default]
org.slf4j#slf4j-api;1.7.30 from central in [default]
org.spark-project.spark#unused;1.0.0 from central in [default]
org.xerial.snappy#snappy-java;1.1.8.2 from central in [default]

```

		modules				artifacts	
conf	number	search	downlded	evicted		number	downlded
default	9	0	0	0		9	0

```

:: retrieving :: org.apache.spark#spark-submit-
parent-b4e4886c-b385-40e0-8941-cde23bf18d8d

```

```

confs: [default]

```

```

0 artifacts copied, 9 already retrieved (0kB/11ms)

```

```

Setting default log level to "WARN".

```

```

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).

```

```

22/04/04 16:44:33 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker

```

```

22/04/04 16:44:33 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster

```

```

22/04/04 16:44:33 INFO org.apache.spark.SparkEnv: Registering
BlockManagerMasterHeartbeat

```

```

22/04/04 16:44:33 INFO org.apache.spark.SparkEnv: Registering
OutputCommitCoordinator

```

```

22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.apache.spark_spark-sql-kafka-0-10_2.12-3.1.2.jar
added multiple times to distributed cache.

```

```

22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.apache.spark_spark-token-provider-
kafka-0-10_2.12-3.1.2.jar added multiple times to distributed cache.

```

```

22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.apache.kafka_kafka-clients-2.6.0.jar added multiple
times to distributed cache.

```

```

22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.apache.commons_commons-pool2-2.6.2.jar added
multiple times to distributed cache.
22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.spark-project.spark_unused-1.0.0.jar added multiple
times to distributed cache.
22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/com.github.luben_zstd-jni-1.4.8-1.jar added multiple
times to distributed cache.
22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.lz4_lz4-java-1.7.1.jar added multiple times to
distributed cache.
22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.xerial.snappy_snappy-java-1.1.8.2.jar added multiple
times to distributed cache.
22/04/04 16:44:36 WARN org.apache.spark.deploy.yarn.Client: Same path resource
file:///root/.ivy2/jars/org.slf4j_slf4j-api-1.7.30.jar added multiple times to
distributed cache.

```

```

[3]: df = spark.readStream.format('kafka').option('kafka.bootstrap.servers', '10.188.
    ↪0.2:9092').option("startingOffsets", "earliest").option('subscribe', '
    ↪'iris-data').option("failOnDataLoss", "false").load()
df = df.selectExpr("CAST(value AS STRING)")

schema = StructType()\
    .add("sepal_length", FloatType())\
    .add("sepal_width", FloatType())\
    .add("petal_length", FloatType())\
    .add("petal_width", FloatType())\
    .add("species", StringType())

print(df.isStreaming)

df.printSchema()

df = df.select(f.from_json(f.decode(df.value, 'utf-8'), schema=schema).
    ↪alias("input"))
df = df.select("input.*")

```

```

True
root
|-- value: string (nullable = true)

```

```

[4]: query3 = df.writeStream.outputMode('update').format('console').start()

```

```

22/04/04 16:44:44 WARN org.apache.spark.sql.streaming.StreamingQueryManager:
Temporary checkpoint location created which is deleted normally when the query

```

didn't fail: /tmp/temporary-e61cb8ee-a68b-490f-b8b2-6e5c31d7beb1. If it's required to delete it under any circumstances, please set spark.sql.streaming.forceDeleteTempCheckpointLocation to true. Important to know deleting temp checkpoint folder is best effort.
 22/04/04 16:44:44 WARN org.apache.spark.sql.streaming.StreamingQueryManager: spark.sql.adaptive.enabled is not supported in streaming DataFrames/Datasets and will be disabled.

```
[5]: model_path = 'gs://big-data-lab-madhurj/Ass7/Pipeline_Model'
model = PipelineModel.load(model_path)
print('Model Loaded....')

predictions = model.transform(df)

mapping = dict(zip([0.0,1.0,2.0],
    ↳ ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']))
mapping_expr = f.create_map([f.lit(x) for x in chain(*mapping.items())])
output_df = predictions.withColumn('prediction', mapping_expr[f.
    ↳ col("prediction")])[:, ['prediction', 'species']]

output_df = output_df.withColumn('correct', f.when((f.
    ↳ col('prediction')== 'Iris-setosa') & (f.col('species')== 'Iris-setosa'),1).
    ↳ when((f.col('prediction')== 'Iris-versicolor') & (f.
    ↳ col('species')== 'Iris-versicolor'),1).when((f.
    ↳ col('prediction')== 'Iris-virginica') & (f.
    ↳ col('species')== 'Iris-virginica'),1).otherwise(0))

df_acc = output_df.select(f.format_number(f.avg('correct')*100,2).
    ↳ alias('accuracy'))

output_df2 = output_df[:, ['prediction', 'species', 'correct']]
output_df2.createOrReplaceTempView('output')
```

22/04/04 16:44:45 WARN org.apache.hadoop.util.concurrent.ExecutorHelper: Thread (Thread[GetFileInfo #1,5,main]) interrupted:
 java.lang.InterruptedException
 at
 com.google.common.util.concurrent.AbstractFuture.get(AbstractFuture.java:510)
 at com.google.common.util.concurrent.FluentFuture\$TrustedFuture.get(FluentFuture.java:88)
 at org.apache.hadoop.util.concurrent.ExecutorHelper.logThrowableFromAfterExecute(ExecutorHelper.java:48)
 at org.apache.hadoop.util.concurrent.HadoopThreadPoolExecutor.afterExecute(HadoopThreadPoolExecutor.java:90)
 at
 java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1157)
 at

```
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:750)
```

```
-----
Batch: 0
-----
```

```
+-----+-----+-----+-----+-----+
|sepal_length|sepal_width|petal_length|petal_width|    species|
+-----+-----+-----+-----+-----+
|         4.9|         3.0|         1.4|         0.2|Iris-setosa|
|         4.7|         3.2|         1.3|         0.2|Iris-setosa|
|         4.6|         3.1|         1.5|         0.2|Iris-setosa|
|         5.0|         3.6|         1.4|         0.2|Iris-setosa|
|         5.4|         3.9|         1.7|         0.4|Iris-setosa|
|         4.6|         3.4|         1.4|         0.3|Iris-setosa|
|         5.0|         3.4|         1.5|         0.2|Iris-setosa|
|         4.4|         2.9|         1.4|         0.2|Iris-setosa|
|         4.9|         3.1|         1.5|         0.1|Iris-setosa|
|         5.4|         3.7|         1.5|         0.2|Iris-setosa|
|         4.8|         3.4|         1.6|         0.2|Iris-setosa|
|         4.8|         3.0|         1.4|         0.1|Iris-setosa|
|         4.3|         3.0|         1.1|         0.1|Iris-setosa|
|         5.8|         4.0|         1.2|         0.2|Iris-setosa|
|         5.7|         4.4|         1.5|         0.4|Iris-setosa|
|         5.4|         3.9|         1.3|         0.4|Iris-setosa|
|         5.1|         3.5|         1.4|         0.3|Iris-setosa|
|         5.7|         3.8|         1.7|         0.3|Iris-setosa|
|         5.1|         3.8|         1.5|         0.3|Iris-setosa|
|         5.4|         3.4|         1.7|         0.2|Iris-setosa|
+-----+-----+-----+-----+-----+
```

```
only showing top 20 rows
```

```
Model Loaded...
```

```
22/04/04 16:44:59 WARN org.apache.spark.ml.feature.StringIndexerModel: Input
column class does not exist during transformation. Skip StringIndexerModel for
this column.
```

```
[ ]: query1 = output_df2.writeStream.queryName("output").outputMode('update').
    ↪format('console').start()
query2 = df_acc.writeStream.outputMode('update').format('console').start()

query1.awaitTermination()
query2.awaitTermination()
```

```
22/04/04 16:44:59 WARN org.apache.spark.sql.streaming.StreamingQueryManager:
Temporary checkpoint location created which is deleted normally when the query
```

didn't fail: /tmp/temporary-8745b3c2-deb6-466c-8cb1-3e71326cd594. If it's required to delete it under any circumstances, please set spark.sql.streaming.forceDeleteTempCheckpointLocation to true. Important to know deleting temp checkpoint folder is best effort.

22/04/04 16:44:59 WARN org.apache.spark.sql.streaming.StreamingQueryManager: spark.sql.adaptive.enabled is not supported in streaming DataFrames/Datasets and will be disabled.

22/04/04 16:44:59 WARN org.apache.spark.sql.streaming.StreamingQueryManager: Temporary checkpoint location created which is deleted normally when the query didn't fail: /tmp/temporary-10ef7f80-0603-4e4e-aa26-71dd181266d4. If it's required to delete it under any circumstances, please set spark.sql.streaming.forceDeleteTempCheckpointLocation to true. Important to know deleting temp checkpoint folder is best effort.

22/04/04 16:44:59 WARN org.apache.spark.sql.streaming.StreamingQueryManager: spark.sql.adaptive.enabled is not supported in streaming DataFrames/Datasets and will be disabled.

Batch: 0

```
+-----+
|accuracy|
+-----+
|   83.89|
+-----+
```

Batch: 0

```
+-----+-----+-----+
| prediction|  species|correct|
+-----+-----+-----+
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
|Iris-setosa|Iris-setosa|    1|
```

Iris-setosa Iris-setosa	1
Iris-setosa Iris-setosa	1
Iris-setosa Iris-setosa	1
Iris-setosa Iris-setosa	1
+-----+-----+-----+	

only showing top 20 rows